



Database fundamentals

Part 3



Agenda

- Limit
- Group by
- Having
- SQL Alias
- Subqueries
- UNION, UNION ALL
- Join
- Views
- Case ...when
- Stored Procedure



LIMIT Clause

The SQL **LIMIT** clause

- To return a **specific number of records** use **limit** clause.
- Useful for pagination or when you don't need the full result.
- The LIMIT clause is useful on large tables with thousands of records.
- Returning a large number of records can impact performance.
- **Note:**
 - Not all database systems support the Limit clause.
 - **MySQL** supports the **LIMIT** clause to select a limited number of records.

LIMIT Clause...

- MySQL Syntax:

```
SELECT column_name(s)  
FROM table_name  
WHERE condition  
LIMIT number;
```

LIMIT...

examples

CustomerID	CustomerName	ContactName	Address	City	PostalCode	Country
1	Alfreds Futterkiste	Maria Anders	Obere Str. 57	Berlin	12209	Germany
2	Ana Trujillo Emparedados y helados	Ana Trujillo	Avda. de la Constitución 2222	México D.F.	05021	Mexico
3	Antonio Moreno Taquería	Antonio Moreno	Mataderos 2312	México D.F.	05023	Mexico
4	Around the Horn	Thomas Hardy	120 Hanover Sq.	London	WA1 1DP	UK

- The following SQL statement **selects the first three records** from the "Customers" table (for MySQL):
 - SELECT * FROM Customers LIMIT 3;**

Result:

Number of Records: 3

CustomerID	CustomerName	ContactName	Address	City	PostalCode	Country
1	Alfreds Futterkiste	Maria Anders	Obere Str. 57	Berlin	12209	Germany
2	Ana Trujillo Emparedados y helados	Ana Trujillo	Avda. de la Constitución 2222	México D.F.	05021	Mexico
3	Antonio Moreno Taquería	Antonio Moreno	Mataderos 2312	México D.F.	05023	Mexico

LIMIT...OFFSET Clause

- **OFFSET:** tells MySQL where to start returning rows.

- **MySQL Syntax:**

```
SELECT column_name(s)
FROM table_name
WHERE condition
LIMIT number
OFFSET number;
```

- **Example:** Skip the first 5 rows, then return the next 5.
- Skips rows 1–5
- Returns rows 6–10

```
SELECT *
FROM employees
LIMIT 5 OFFSET 5;
```



Group By and Having Clauses

The SQL **Group by** clause

- For use with aggregate functions (min, max, count, sum, avg).
- multiple values returned from SQL query with aggregate function (via GROUP BY)
- **Example:** Count employees in each department

```
SELECT department, COUNT(*) AS total_employees  
FROM employees  
GROUP BY department;
```

The SQL **Having** clause

- For use with GROUP BY.
- Like WHERE clause, but used to filter groups after aggregation, not on individual rows.
- You **can't** use **WHERE** with aggregate functions, that's why we use **HAVING**.
- **Example:** Show departments with more than 5 employees

```
SELECT department, COUNT(*) AS total_employees
FROM employees
GROUP BY department
HAVING COUNT(*) > 5;
```

- Explanation:
 - First groups employees by department.
 - Then filters only groups where COUNT(*) > 5.

Difference Between **WHERE** and **HAVING**

- **WHERE**: filters rows before grouping (on raw data).
- **HAVING**: filters groups after aggregation (on aggregated results).
- **Example** combining both:

```
SELECT department, COUNT(*) AS total_employees
FROM employees
WHERE salary > 3000      -- filter rows first
GROUP BY department
HAVING COUNT(*) > 2;     -- then filter groups
```



SQL *Aliases*

SQL Aliases

- SQL aliases are used to give a **table**, or a **column** in a table, a **temporary name**.
- Aliases are often used **to make column names more readable**.
- An alias **only exists for the duration of that query**.
- An alias is created with the **AS** keyword.
- **Alias Column Syntax**
 - `SELECT column_name AS alias_name
FROM table_name;`
- **Alias Table Syntax**
 - `SELECT column_name(s)
FROM table_name AS alias_name;`

SQL Aliases example

- The following SQL statement creates two aliases, one for the CustomerID column and one for the CustomerName column:
 - **SELECT** CustomerID **AS** **ID**, CustomerName **AS** **Customer**
FROM Customers;

Result:

Number of Records: 91

ID	Customer
1	Alfreds Futterkiste
2	Ana Trujillo Emparedados y helados
3	Antonio Moreno Taquería
4	Around the Horn
5	Berglunds snabbköp
6	Blauer See Delikatessen
7	Blondel père et fils

SQL Aliases example

- The following SQL statement creates two aliases, one for the CustomerName column and one for the ContactName column.
- **Note:** It requires **double quotation** marks or **square brackets** if the **alias name contains spaces**
 - **SELECT** CustomerName **AS** Customer, ContactName **AS** [Contact Person]
FROM Customers;

Result:

Number of Records: 91

Customer	Contact Person
Alfreds Futterkiste	Maria Anders
Ana Trujillo Emparedados y helados	Ana Trujillo
Antonio Moreno Taquería	Antonio Moreno
Around the Horn	Thomas Hardy

SQL Aliases example

- The following SQL statement creates an alias named "Address" that combine four columns (Address, PostalCode, City and Country):
 - `SELECT CustomerName, concat(Address , ' , ' , PostalCode , ' ' , City , ' , ' , Country) AS Address
FROM Customers;`

Result:

Number of Records: 91

CustomerName	Address
Alfreds Futterkiste	Obere Str. 57, 12209 Berlin, Germany
Ana Trujillo Emparedados y helados	Avda. de la Constitución 2222, 05021 México D.F., Mexico
Antonio Moreno Taquería	Mataderos 2312, 05023 México D.F., Mexico
Around the Horn	120 Hanover Sq., WA1 1DP London, UK
Berglunds snabbköp	Berguvsvägen 8, S-958 22 Luleå, Sweden

SQL Aliases

- Aliases can be useful when:
 - There are more than one table involved in a query
 - Functions are used in the query
 - Column names are big or not very readable
 - Two or more columns are combined together

Subqueries

The SQL Subquery

- A subquery is a query inside another query.
- It is enclosed in parentheses () and can be used in SELECT, INSERT, UPDATE, or DELETE statements.
- Think of it like:

First run the inner query, then use its result in the outer query.

The SQL Subquery examples

- **Example:** Find employees who earn more than the average salary.
 - The inner query: `SELECT AVG(salary) FROM employees` → returns the average salary.
 - The outer query: selects employees with salary above that value.

```
SELECT first_name, last_name, salary
FROM employees
WHERE salary > (
    SELECT AVG(salary)
    FROM employees
);
```

- **Example:** Find employees who work in departments that are located in New York.

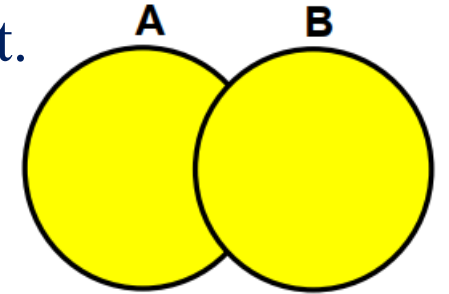
```
SELECT first_name, last_name
FROM employees
WHERE department IN (
    SELECT department_id
    FROM departments
    WHERE location = 'New York'
);
```



UNION and UNION ALL Clauses

SQL UNION...

- Combines the results of two or more SELECT queries into one result set.
- Removes duplicate rows by default.
- The expressions in the SELECT lists must match in number.
- The data type of each column in the second query must match the data type of its corresponding column in the first query.
- **Example:** Display the current and previous job details of all employees.
- Display each employee only once.



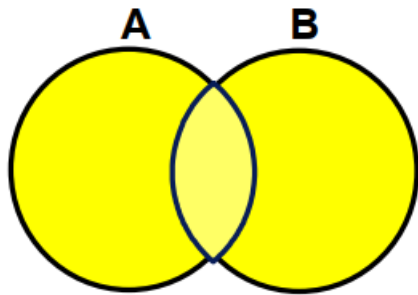
```
SELECT employee_id, job_id  
FROM employees
```

```
UNION
```

```
SELECT employee_id, job_id  
FROM job_history;
```

SQL UNION All...

- Same as **UNION**, but it **does not** remove duplicates.
- Keeps all rows, even if repeated.



- If You Use UNION ALL This will keep duplicates.

```
SELECT employee_id, job_id
FROM employees

UNION ALL

SELECT employee_id, job_id
FROM job_history;
```

◆ Example with Different Tables

Suppose we have two tables:

- `customers(id, name, email)`
- `suppliers(id, name, email)`

👉 Get a list of all people (customers + suppliers):

sql

```
SELECT name, email FROM customers
UNION
SELECT name, email FROM suppliers;
```

👉 Get them with duplicates included:

sql

```
SELECT name, email FROM customers
UNION ALL
SELECT name, email FROM suppliers;
```



SQL JOIN

Types of Joins with Examples

SQL JOIN


- Obtaining Data from Multiple Tables.
- Write SELECT statements to access data from more than one table using join.

EMPLOYEES

	EMPLOYEE_ID	LAST_NAME	DEPARTMENT_ID
1	100	King	90
2	101	Kochhar	90
3	102	De Haan	90
...			
18	202	Fay	20
19	205	Higgins	110
20	206	Gietz	110

DEPARTMENTS

	DEPARTMENT_ID	DEPARTMENT_NAME	LOCATION_ID
1	10	Administration	1700
2	20	Marketing	1800
3	50	Shipping	1500
4	60	IT	1400
5	80	Sales	2500
6	90	Executive	1700
7	110	Accounting	1700
8	190	Contracting	1700



	EMPLOYEE_ID	DEPARTMENT_ID	DEPARTMENT_NAME
1	200	10	Administration
2	201	20	Marketing
3	202	20	Marketing
4	124	50	Shipping
5	144	50	Shipping

...			
18	205	110	Accounting
19	206	110	Accounting

SQL JOINS

- A JOIN clause is used to combine rows from two or more tables, based on a related column between them (usually a primary key / foreign key relationship).
 - Let's look at a selection from the "Orders" table:

OrderID	CustomerID	OrderDate
10308	2	1996-09-18
10309	37	1996-09-19
10310	77	1996-09-20

- Then look at a selection from the "Customers" table:

CustomerID	CustomerName	ContactName	Country
1	Alfreds Futterkiste	Maria Anders	Germany
2	Ana Trujillo Emparedados y helados	Ana Trujillo	Mexico
3	Antonio Moreno Taquería	Antonio Moreno	Mexico

SQL JOINS...

- Notice that the "CustomerID" column in the "Orders" table refers to the "CustomerID" in the "Customers" table. The relationship between the two tables above is the "CustomerID" column.
- Then, we can create the following SQL statement (that contains a JOIN), that **selects records that have matching values in both tables**:
 - **SELECT** Orders.OrderID, Customers.CustomerName, Orders.OrderDate
FROM Orders
JOIN Customers **ON** Orders.CustomerID=Customers.CustomerID;

OrderID	CustomerName	OrderDate
10308	Ana Trujillo Emparedados y helados	9/18/1996
10365	Antonio Moreno Taquería	11/27/1996
10383	Around the Horn	12/16/1996
10355	Around the Horn	11/15/1996
10278	Berglunds snabbköp	8/12/1996



Types of JOIN

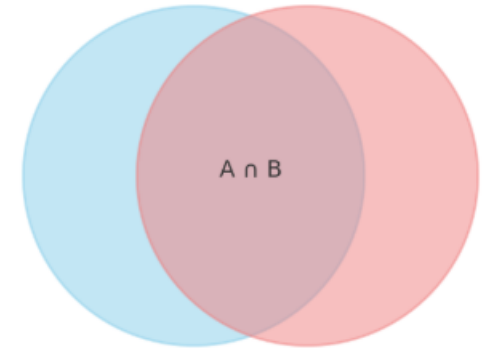


INNER JOIN

- Returns rows with matching values in both tables.
- **Example:** select employee with its department name.
- Shows only employees that belong to a department.

```
SELECT e.employee_id, e.first_name, d.department_name
FROM employees e
INNER JOIN departments d
    ON e.department_id = d.department_id;
```

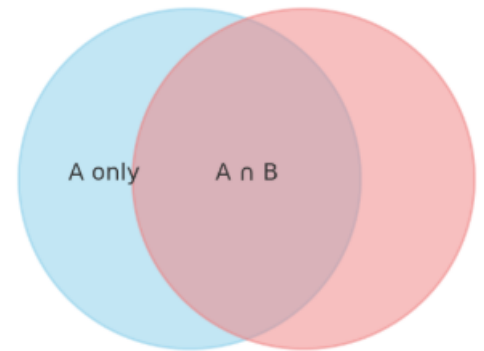
INNER JOIN



LEFT JOIN

- Returns all rows from the left table + matching rows from right.
- If no match, right table columns are NULL appear.
- **Example:**
- Shows all employees, even those not assigned to any department.

LEFT JOIN



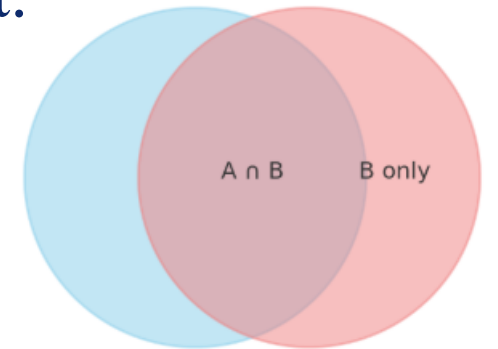
```
SELECT e.employee_id, e.first_name, d.department_name
FROM employees e
LEFT JOIN departments d
    ON e.department_id = d.department_id;
```

RIGHT JOIN

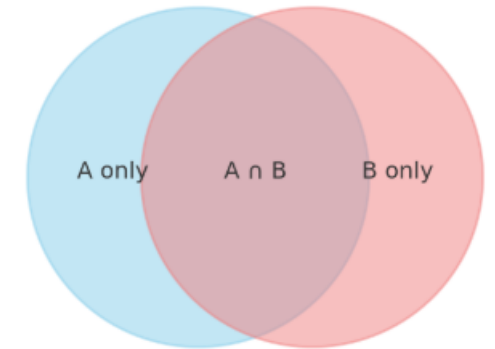
- Returns all rows from the right table + matching rows from left.
- If no match, right table columns are NULL appear.
- **Example:**
- Shows all departments, even if no employees work in them.

```
SELECT e.employee_id, e.first_name, d.department_name
FROM employees e
RIGHT JOIN departments d
    ON e.department_id = d.department_id;
```

RIGHT JOIN



FULL JOIN (Simulated in MySQL)



- Returns all rows from both tables.
- MySQL **does not support** FULL JOIN directly.
- Can be simulated with **UNION** of LEFT JOIN and RIGHT JOIN.

- **Example:**

```
SELECT e.employee_id, e.first_name, d.department_name
FROM employees e
LEFT JOIN departments d
    ON e.department_id = d.department_id

UNION

SELECT e.employee_id, e.first_name, d.department_name
FROM employees e
RIGHT JOIN departments d
    ON e.department_id = d.department_id;
```


SELF JOIN

- A table joined with **itself** (useful for hierarchical data, like managers and employees).
- It's often used to find relationships inside the same table.
- **Example:**
- Shows each employee with their manager's name.

```
SELECT e.employee_id, e.first_name, m.first_name AS manager_name
FROM employees e
JOIN employees m
    ON e.manager_id = m.employee_id;
```

ON Keyword

- Used in the JOIN clause to define join conditions (how two tables are linked).
- Used with all types of joins (INNER JOIN, LEFT JOIN, RIGHT JOIN, etc.).
- It defines the relationship/condition between the two tables.
- Keeps join condition separate from filtering conditions.
- If the condition in ON is not satisfied, the row will not appear (except in outer joins where NULLs may appear).
- Without ON, SQL wouldn't know which columns to use to connect the tables.

JOIN Summery

- INNER JOIN → matching rows only.
- LEFT JOIN → all left + matches from right.
- RIGHT JOIN → all right + matches from left.
- FULL JOIN → all rows from both (MySQL: simulate with UNION)
- SELF JOIN → table joined to itself

VIEWS

The SQL VIEWS

- Views provide users-controlled access to tables.
- Base Table—table containing the raw data.
- Dynamic View
 - A “virtual table” created dynamically upon request by a user.
 - No data actually stored; instead, data from base table made available to user.
 - Based on SQL SELECT statement on base tables or other views.
- Materialized View
 - Copy or replication of data.
 - Data actually stored.
 - Must be refreshed periodically to match corresponding base tables.

The SQL **VIEWS** example

- View has a name.
- View is based on a SELECT statement.
- **Create view**

```
-- Create a view to show employee full names and their departments
CREATE VIEW employee_summary AS
SELECT id,
       CONCAT(first_name, ' ', last_name) AS full_name,
       department
FROM employees;
```

- Now you can query it like a table.

```
SELECT * FROM employee_summary;
```

The SQL VIEWS

- **Modify View**

```
ALTER VIEW employee_summary AS
SELECT id,
       CONCAT(first_name, ' ', last_name) AS full_name,
       department,
       salary
FROM employees;
```

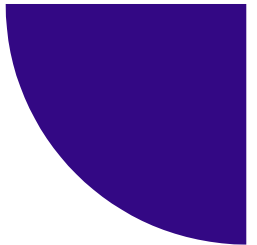
- **Drop view**

```
-- Drop view view_name;

Drop view employee_summary;
```

Advantages of Views

- To restrict data access.
- To make complex queries easy.
- To provide data independence.
- To present different views of the same data.



Case ...when

CASE...WHEN

- The CASE WHEN statement in MySQL is used to add conditional logic inside SQL queries — like an IF...ELSE in programming.
- Syntax:

```
CASE
  WHEN condition1 THEN result1
  WHEN condition2 THEN result2
  ...
  ELSE default_result
END
```

Example:

```
SELECT first_name,
       salary,
       CASE
         WHEN salary > 10000 THEN 'High'
         WHEN salary BETWEEN 5000 AND 10000 THEN 'Medium'
         ELSE 'Low'
       END AS salary_level
FROM employees;
```



Stored Procedure



Stored Procedure

- A Stored Procedure is a set of SQL statements that you save in the database and run later by calling it.
- It's like a function in programming: reusable, parameterized, and stored on the server.
- Encapsulates complex queries.
- Improves performance (stored on server).
- Reusable by multiple apps.
- Can include loops, conditions, variables.

Stored Procedure

- Creating Syntax

```
DELIMITER $$

CREATE PROCEDURE procedure_name()
BEGIN
    -- SQL statements
END$$

DELIMITER ;
```

Using or calling syntax

```
-- ♦ Basic call (no parameters)
CALL procedure_name();

-- ♦ With one IN parameter (just pass a value)
CALL procedure_name(123);

-- ♦ With IN and OUT parameter
-- IN → pass a value
-- OUT → capture result in a user variable
CALL procedure_name(123, @out_value);

-- ♦ With multiple parameters (mix of IN, OUT, INOUT possible)
CALL procedure_name('Alice', 5000, @status);
```

- **DELIMITER \$\$** → temporarily changes the statement terminator so MySQL doesn't stop at the first ;.
- **CREATE PROCEDURE** → defines the procedure.
- **BEGIN ... END** → block of SQL code.

Stored Procedure examples

- Example 1: Simple Procedure

```
DELIMITER $$

CREATE PROCEDURE GetAllEmployees()
BEGIN
    SELECT * FROM employees;
END$$

DELIMITER ;
```

- Calling it

```
CALL GetAllEmployees();
```

Stored Procedure examples

- Example 2: With Input Parameter

```
DELIMITER $$

CREATE PROCEDURE GetEmployeesByDept(IN dep_id INT)
BEGIN
    SELECT employee_id, first_name, salary
    FROM employees
    WHERE department_id = dep_id;
END$$

DELIMITER ;
```

- Calling it

```
CALL GetEmployeesByDept(10);
```

Stored Procedure examples

- Example 3: With Output Parameter

```
DELIMITER $$

CREATE PROCEDURE CountEmployeesByDept(IN dep_id INT, OUT emp_count INT)
BEGIN
    SELECT COUNT(*) INTO emp_count
    FROM employees
    WHERE department_id = dep_id;
END$$

DELIMITER ;
```

- Calling it

```
CALL CountEmployeesByDept(10, @total);
SELECT @total; -- shows number of employees in dept 10
```


Stored Procedure examples

- Example 4: Procedure with IF / CASE

Calling it

```
DELIMITER $$

CREATE PROCEDURE SalaryLevel(IN emp_id INT, OUT level VARCHAR(10))
BEGIN
    DECLARE emp_salary DECIMAL(10,2);

    SELECT salary INTO emp_salary
    FROM employees
    WHERE employee_id = emp_id;

    CASE
        WHEN emp_salary > 10000 THEN SET level = 'High';
        WHEN emp_salary BETWEEN 5000 AND 10000 THEN SET level = 'Medium';
        ELSE SET level = 'Low';
    END CASE;
END$$

DELIMITER ;
```

```
-- 1 Initialize the OUT variable (to store the procedure's result)
SET @emp_level = '';

-- 2 Call the stored procedure with:
--     - IN parameter: employee_id = 101
--     - OUT parameter: @emp_level (to capture the level result)
CALL SalaryLevel(101, @emp_level);

-- 3 Display the value of the OUT parameter returned by the procedure
SELECT @emp_level;
```

Any Questions ?